




Parallel and Distributed Computing

Dr. Ali Sayyed
Department of Computer Science
National University of Computer & Emerging Sciences



Shared Memory Programming Models

OpenMP



Programming Multicore Systems



- Most of the **common languages** were created prior to the multicore age.
- As a result, many languages **cannot implicitly** (or **automatically**) employ multicore processors to speed up the execution of a program.
- Instead, **programmers** must specifically write software to leverage the multiple cores on a system.

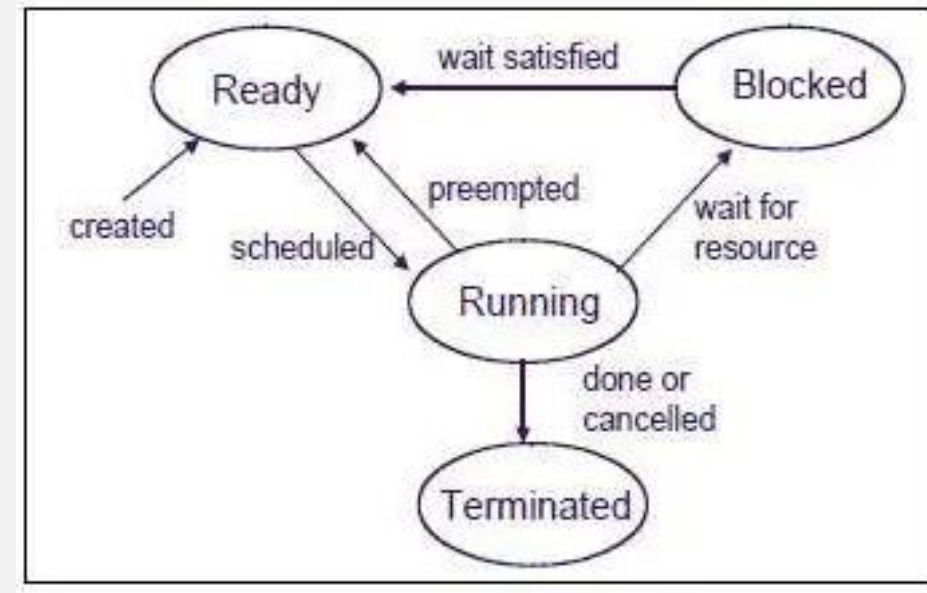


Expediting Process Execution

- One way to speed up the execution of a single process is to decompose it into lightweight, independent execution flows (a **unit of work**) called **threads**.
- The OS **schedules threads** in the same manner as it schedules processes.
- On a multicore processor, the OS can speed up the execution of a multithreaded program by **scheduling** different threads to run on **separate cores**.
- The **maximum number of threads** that can execute in parallel is equal to the number of physical cores on the system.
 - The excess threads must **wait** their turn to execute
- Threads communicate **implicitly** by writing/ reading shared variables.

Thread Lifecycle

- A thread is “born” by the **create()** function, which places it in the ready state.
- A thread **runs when it is scheduled**.
- The running thread may be **blocked because it must wait** for some resource
- It may be **preempted** either because a higher **priority** thread is ready to run, or its **time slice** has expired.
- When the **resource becomes available** a blocked thread transitions to the Ready state and will eventually be scheduled to run again.
- Finally, a thread is **terminated** when it is done, or another thread requests its cancellation.





Shared Memory Programming



Several Thread Libraries/systems

- **PTHREADS is the POSIX Standard**
 - Relatively low level, possibly slow and relatively heavyweight
 - IEEE
- **OpenMP** standard for application-level programming
 - Support for scientific programming on shared memory
 - openmp.org
- **TBB: Thread Building Blocks**
 - Intel
- **Java threads**
 - Built on top of POSIX threads
 - Object within Java language



OpenMP Overview




- OpenMP (Open Multi-Processing) is a programming API that simplifies writing multi-threaded, shared-memory parallel programs in C, C++, and Fortran.
- **Compiler Directives:** OpenMP uses compiler directives (pragmas) to specify parallel regions of code, which the compiler then translates into threaded code
- **Runtime Library Routines:** OpenMP provides a runtime library that allows developers to manage threads, control the execution environment, and manipulate locks on memory locations.
- **Environment Variables:** OpenMP also uses environment variables to influence the runtime behavior of parallel programs.





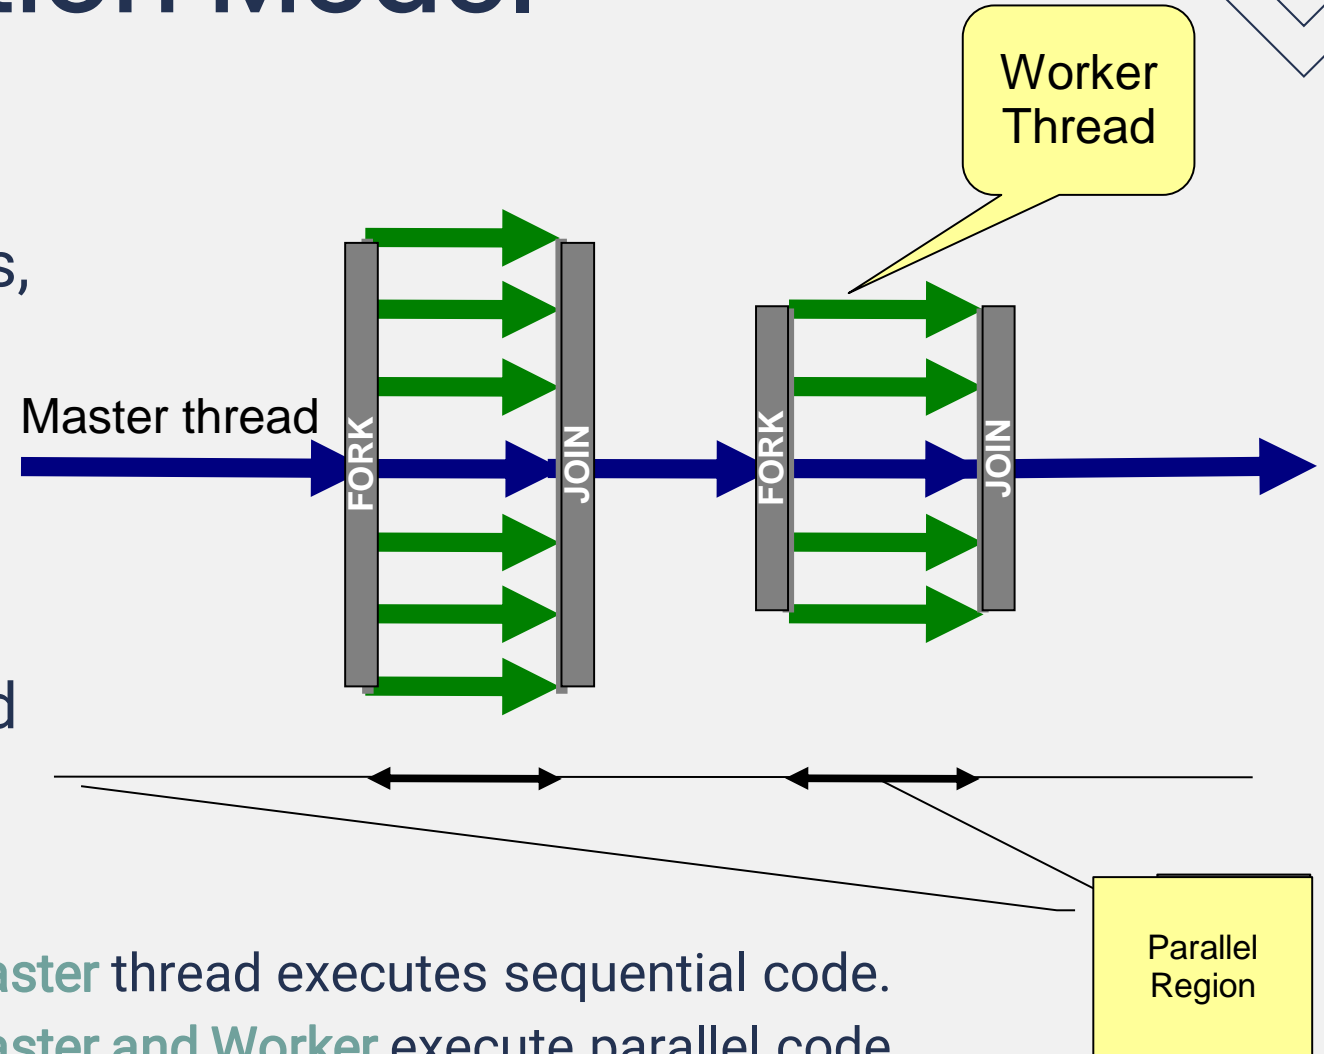
OpenMP API Overview

- Latest Specification is OpenMP 6.0 released in Nov 2024
- Learning more about MP
 - <https://hpc.llnl.gov/tuts/openMP/>
 - <https://www.openmp.org>
 - <http://tinyurl.com/OpenMP-Tutorial>
- A number of Compilers and Tools from various vendors and open source community initiatives implement the OpenMP API.
 - <https://www.openmp.org/resources/openmp-compilers-tools/>

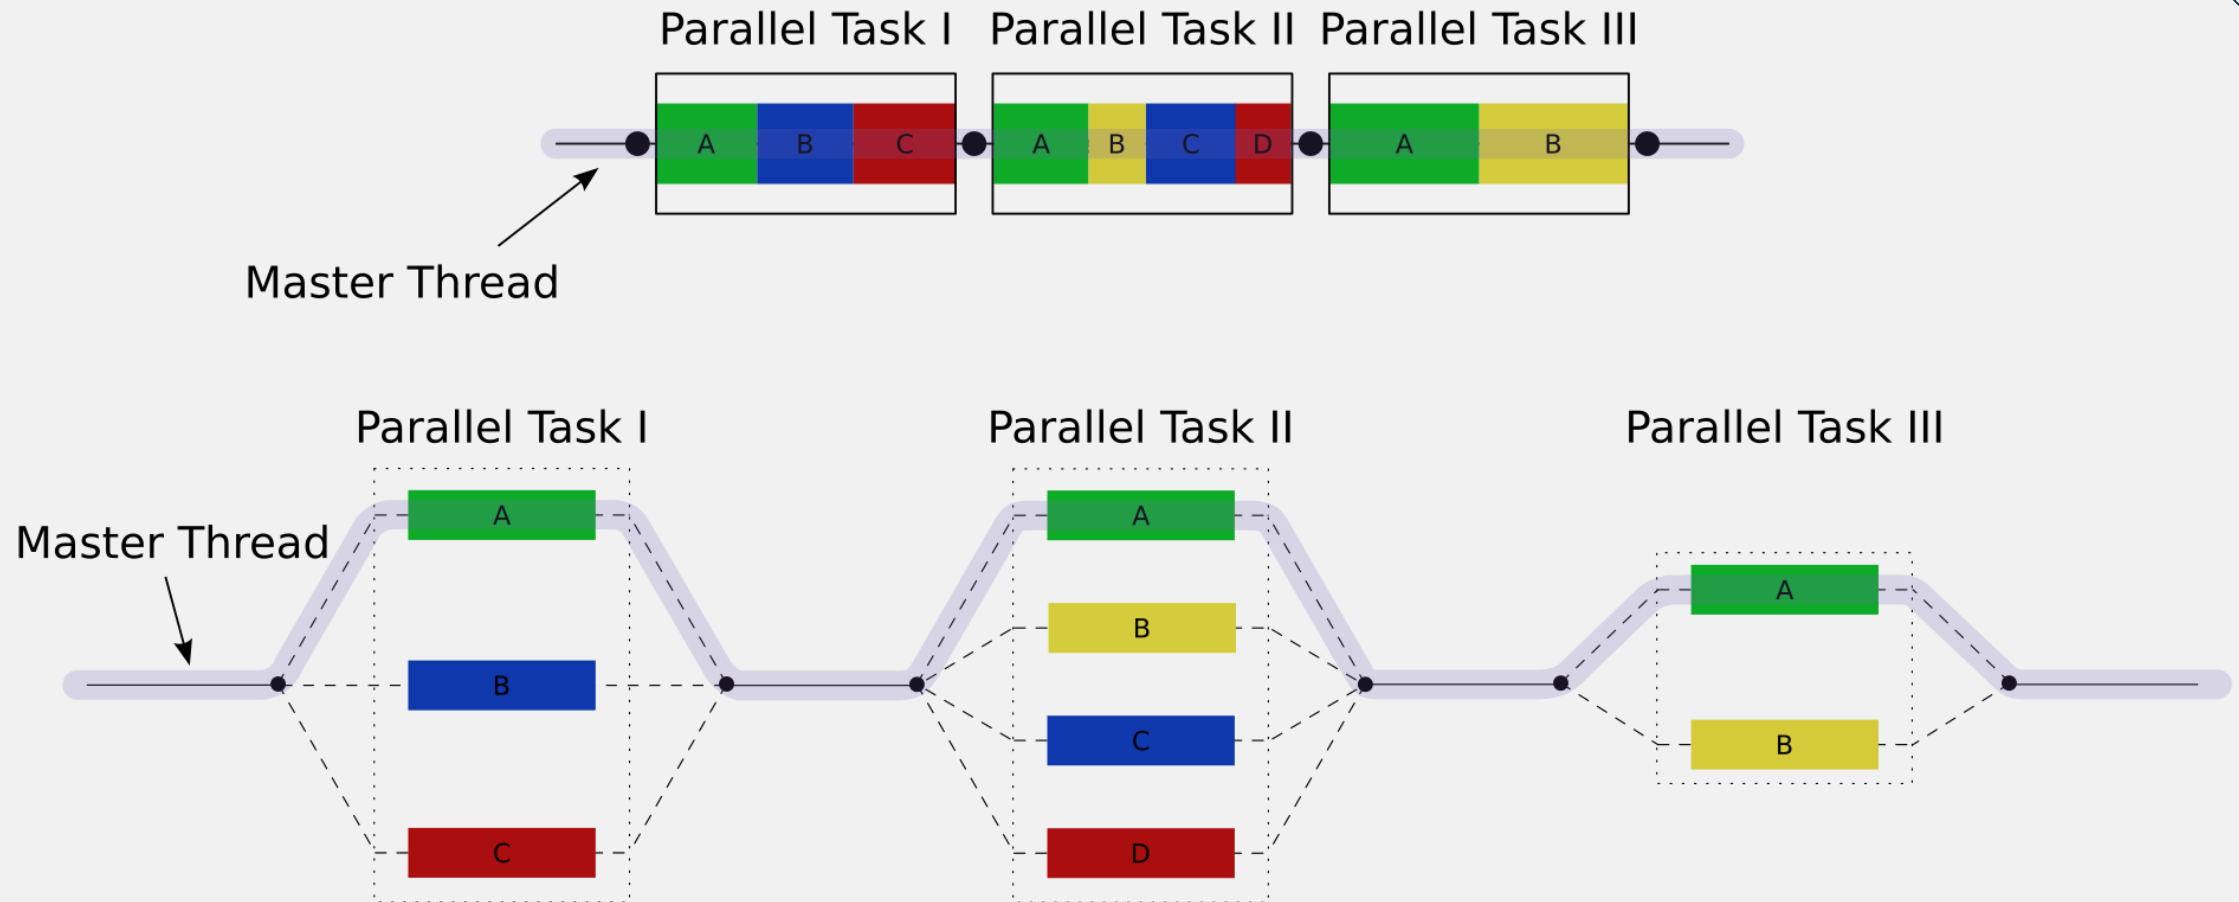


Execution Model

- OpenMP program starts single threaded
- To create additional threads, user starts a parallel region
 - additional threads are launched to create a team
 - original (master) thread is part of the team
- Repeat parallel regions as necessary
 - Fork-join model



Fork-Join Model





OpenMP API Example

Sequential code:

```
statement1;  
statement2;  
statement3;
```

Assume we want to execute statement 2 in parallel, and statement 1 and 3 sequentially.





OpenMP API Example

OpenMP parallel code:

```
statement 1;
```

```
#pragma <specific OpenMP directive>
```

```
statement2;
```

```
statement3;
```

Statement 2 (may be) executed in parallel.

Statement 1 and 3 are executed sequentially.


OpenMP Specification

<https://www.openmp.org/spec-html/5.0/openmp.html>




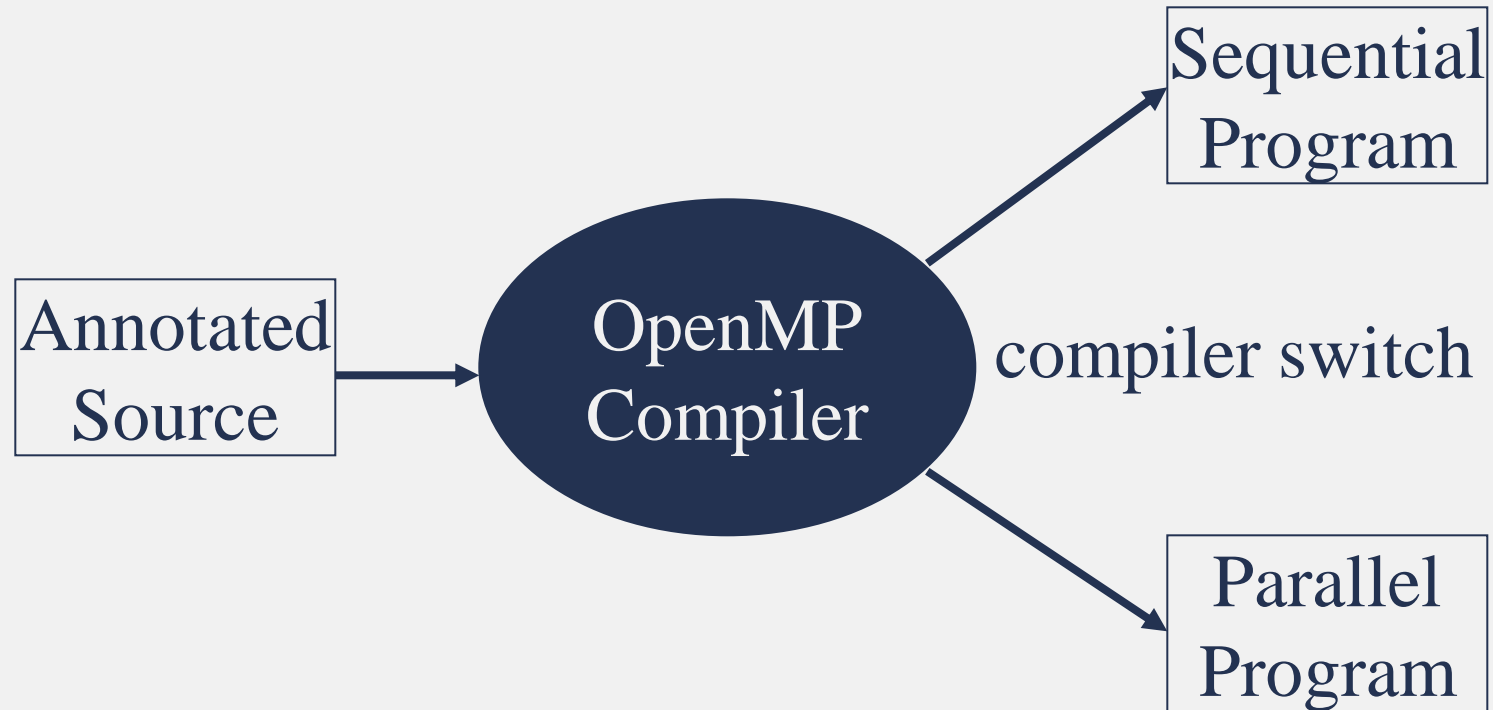
Important Note

- By giving a parallel directive, the user asserts that the program will remain correct if the statement is executed in parallel.
- OpenMP compiler does not check correctness.
- Some tools exist for helping with that.






OpenMP Example Usage





Compiler Switch

- Compile with `-fopenmp` to enable OpenMP if using GNU Compilor
- The Microsoft Visual C/C++ compiler supports the OpenMP 2.0 standard with the `-openmp` switch.
- Compile with `-mp` to enable OpenMP for multicore CPUs on all platforms, if using NVIDIA HPC compiler.
- Compile with `-Qopenmp` on Windows, or just `-qopenmp` or `-fiopenmp` on Linux or Mac OSX, if using Intel Compiler.





OpenMP Example Usage




- If you give sequential switch,
 - comments and pragmas are ignored.
- If you give parallel switch,
 - comments and/or pragmas are read, and
 - cause translation into parallel program.
- Ideally, one source for both sequential and parallel program (**big maintenance plus**).



Tools

- **HPE's** Code Parallelization Assistant, a part of the HPE Cray Programming Environment, can be used to identify and exploit parallelism. This tool helps highlight dependencies or bottlenecks during the optimization phase of program development.
- **Intel VTune Profiler** is a low-overhead performance profiling and analysis tool which helps to find and fix performance bottlenecks quickly.
- **Intel Advisor** is a design and analysis tool to help ensure your applications realize full performance potential on modern Intel processors.
- **Intel Inspector** is an easy-to-use memory and threading error debugger for C, C++, and Fortran applications that run on Windows* and Linux
- **NVIDIA Nsight Systems** for Linux is capable of capturing information about OpenMP events.



OpenMP Directives

- **Parallelization** directives:
 - parallel region
 - parallel for
 - etc.
- **Data environment** directives:
 - shared, private, threadprivate, reduction,
 - etc.
- **Synchronization** directives:
 - barrier, critical
 - etc.



General Rules about Directives

- They always apply to the next statement, which must be a structured block.
- Examples
 - `#pragma omp construct [clause ...]`
`statement`
 - `#pragma omp construct [clause ...]`
`{ statement1; statement2; statement3; }`



OpenMP Parallel Region

- Defines a parallel region, which is code that will be executed by multiple threads in parallel.

```
#pragma omp parallel [clauses]
{
    code_block
}
```

- (Optional) Zero or more clauses.
- Each thread executes the **same code**.
- Each thread waits at the end.